

# The Power of Scale for Parameter-Efficient Prompt Tuning

**Brian Lester\* Rami Al-Rfou Noah Constant**

Google Research

`{brianlester, rmyeid, nconstant}@google.com`

# Abstract

---

- Prompt Tuning
  - In Downstream Task, Learning “Soft Prompt” + Frozen Language Models
  - more competitive with scale
  - benefits in robustness to domain transfer
  - efficient prompt ensembling

# Introduction

---

- How to adapt general-purpose models to downstream tasks?
  - Model tuning (fine-tuning)
    - all model parameters are tuned during adaptation
  - Prompt design (priming)
    - prompts: a task description and/or several examples
    - freezing pre-trained models
    - make a single generalist model simultaneously serve many different tasks

# Introduction

---

- Key drawbacks of prompt-based adaptation
  - Task description is error-prone and requires human involvement
  - The effectiveness depends on the prompt quality
  - Downstream task quality is lower than tuned models
- Solution: Prompt Tuning!

# Prompt Tuning

- Conditional Text Generation of T5 (Classification)

$$\Pr_{\theta}(Y|X)$$

- $X$  is a series of tokens,  $Y$  is a sequence of tokens that represent a class label
- $\theta$  is the weights of the transformers that make up its encoder and decoder

# Prompt Tuning

- Prompting is the approach of adding extra information

$$\Pr_{\theta}(Y|[P; X])$$

- $P$  is a series of tokens prepended (prompts)
  - $P = \{p_1, p_2, \dots, p_n\}$
- The model maximizes the likelihood of the correct  $Y$ , while keeping the model parameters  $\theta$  (fixed)

# Prompt Tuning

---

- Prompt Design
  - selecting prompt tokens from a fixed vocabulary of frozen embeddings
- Prompt Tuning
  - using a fixed prompt of special tokens,  
where only the embeddings of these prompt tokens **can be updated**

# Prompt Tuning

- New conditional generation model

$$\Pr_{\theta; \theta_P}(Y | [P; X])$$

- The model maximizes the likelihood of  $Y$  via backpropagation
- Only applying gradient updates to  $\theta_P$



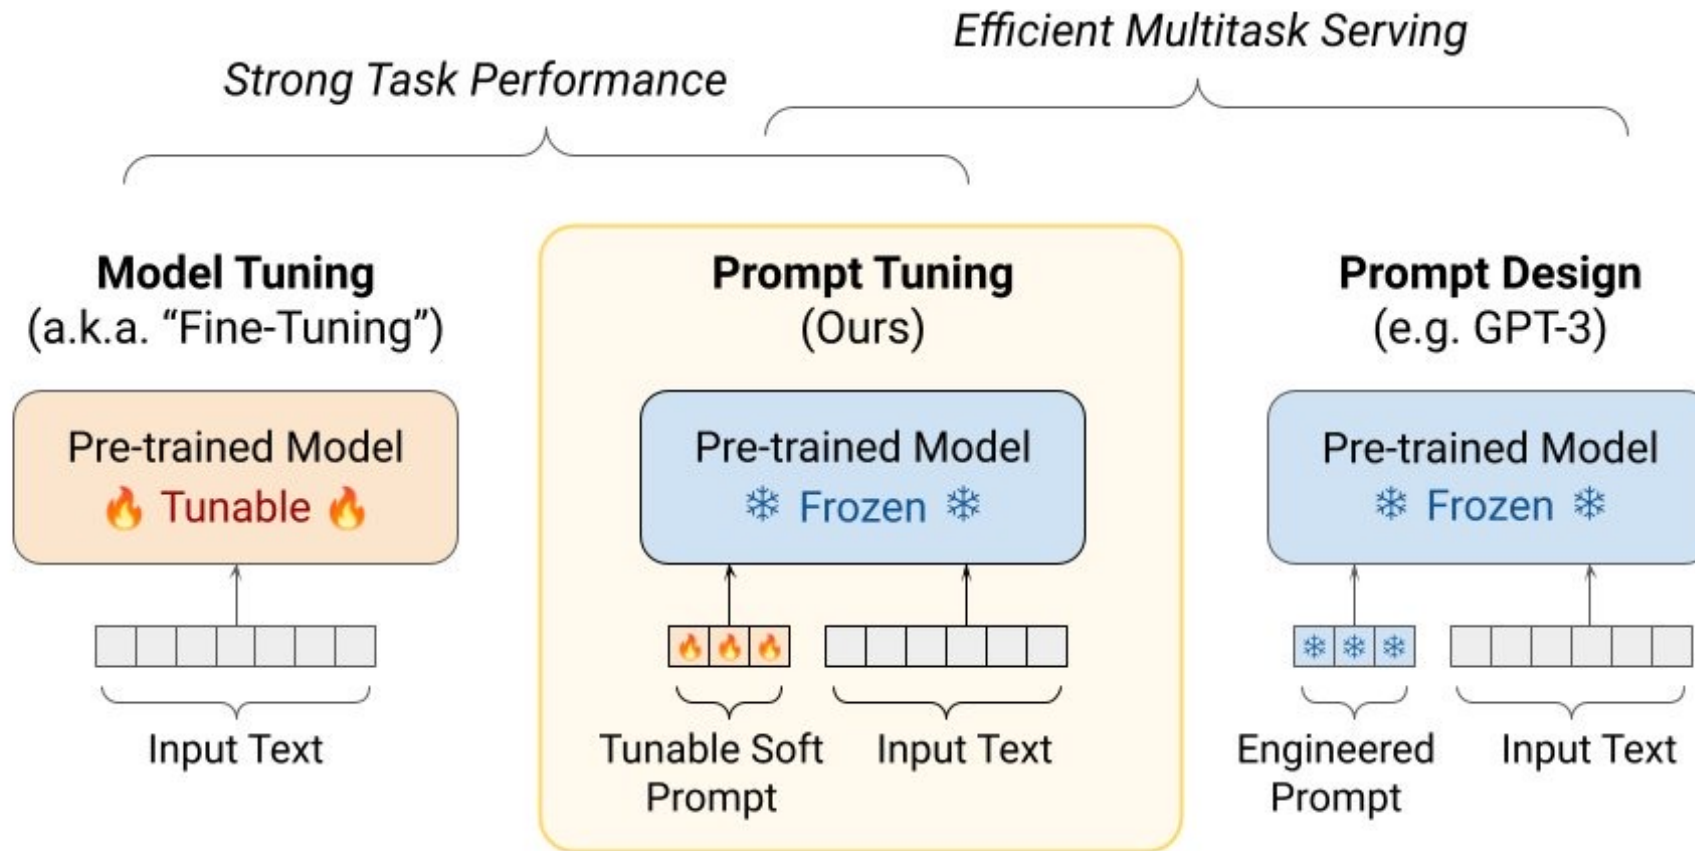
# Prompt Tuning

- New conditional generation model

$$\Pr_{\theta; \theta_P}(Y | [P; X])$$

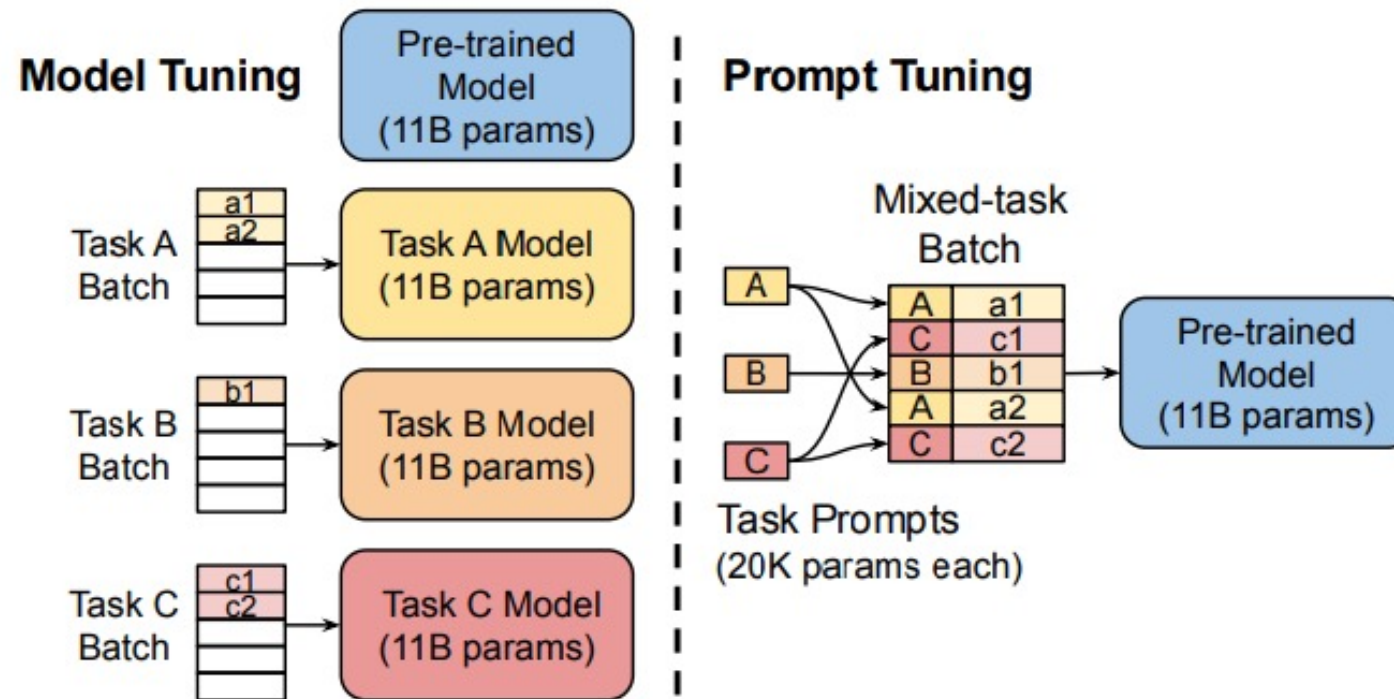
- Given a series of  $n$  tokens,  $\{x_1, x_2, \dots, x_n\}$ ,  
T5 embed the tokens forming a matrix  $X_e \in \mathbb{R}^{n \times e}$ 
  - $e$  is the dimension of the embedding space
- soft-prompts are represented as a parameter  $P_e \in \mathbb{R}^{p \times e}$
- Then, our prompt is  $[P_e; X_e] \in \mathbb{R}^{(p+n) \times e}$ 
  - only the prompt parameters  $P_e$  are updated

# Prompt Tuning



Reference: Google AI

# Prompt Tuning



# Prompt Tuning

- Design Decisions
  - How to initialize the prompt representations?
    - random initialization
    - initialize each prompt token to an embedding drawn from the model's vocab
    - For classification tasks,  
initialize the prompt with embeddings that enumerate the output classes
  - The length of the prompt
    - aim to find a minimal length that still performs well

# Prompt Tuning


- Unlearning Span Corruption

- T5 is tasked with “reconstructing” masked spans in the input text
- This setup is not a good fit for prompt tuning
- Example)
  - Text: “Thank you for inviting me to your party last week”
  - Input: “Thank you <X> me to your party <Y> week”
  - Output: “<X> for inviting <Y> last <Z>”

# Prompt Tuning

- Unlearning Span Corruption
  - Span Corruption
    - just use pre-trained T5 off-the-shelf as our frozen model
  - Span Corruption + Sentinel
    - use the same model, but prepend all downstream targets with a sentinel
  - LM Adaptation
    - continue T5's self-supervised training for a small number of additional steps
    - natural text prefix -> natural text continuation
    - quickly transform T5 into a model more similar to GPT-3


# Results




- Default configuration
  - LM-adapted version of T5
  - initialize using class labels
  - use a prompt length of 100 tokens
- SuperGLUE benchmark
- Prompt Tuning matches the baselines

# Results

- Ablation Study



(a) Prompt length




(b) Prompt initialization




# Results

- Ablation Study




(c) Pre-training method



(d) LM adaptation steps

# Comparison to Similar Approaches



- Prompt Tuning is the most parameter-efficient

# Resilience to Domain Shift

- Prompt tuning prevents the model from modifying its general understanding
- This suggests that prompt tuning may improve robustness to domain shifts
- Zero-shot domain transfer: QA, paraphrase detection

Dataset	Domain	Model	Prompt	$\Delta$
SQuAD	Wiki	94.9 $\pm$ 0.2	94.8 $\pm$ 0.1	-0.1
TextbookQA	Book	54.3 $\pm$ 3.7	<b>66.8</b> $\pm$ 2.9	+12.5
BioASQ	Bio	77.9 $\pm$ 0.4	<b>79.1</b> $\pm$ 0.3	+1.2
RACE	Exam	59.8 $\pm$ 0.6	<b>60.7</b> $\pm$ 0.5	+0.9
RE	Wiki	88.4 $\pm$ 0.1	<b>88.8</b> $\pm$ 0.2	+0.4
DuoRC	Movie	<b>68.9</b> $\pm$ 0.7	67.7 $\pm$ 1.1	-1.2
DROP	Wiki	<b>68.9</b> $\pm$ 1.7	67.1 $\pm$ 1.9	-1.8

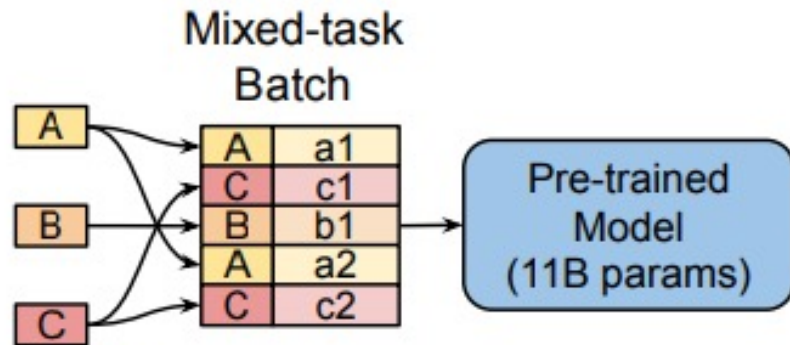
Table 1: F1 mean and stddev for models trained on SQuAD and evaluated on out-of-domain datasets from the MRQA 2019 shared task. Prompt tuning tends to give stronger zero-shot performance than model tuning, especially on datasets with large domain shifts like TextbookQA.

Train	Eval	Tuning	Accuracy	F1
QQP	MRPC	Model	73.1 $\pm$ 0.9	81.2 $\pm$ 2.1
		Prompt	<b>76.3</b> $\pm$ 0.1	<b>84.3</b> $\pm$ 0.3
MRPC	QQP	Model	74.9 $\pm$ 1.3	<b>70.9</b> $\pm$ 1.2
		Prompt	<b>75.4</b> $\pm$ 0.8	69.7 $\pm$ 0.3

Table 2: Mean and stddev of zero-shot domain transfer between two paraphrase detection tasks.

# Prompt Ensembling

- By training N prompts on the same tasks, we create N separate “models” for a task, while still sharing the core language modeling parameters throughout



Dataset	Metric	Average	Best	Ensemble
BoolQ	acc.	91.1	91.3	<b>91.7</b>
CB	acc./F1	99.3 / 99.0	100.00 / 100.00	<b>100.0 / 100.0</b>
COPA	acc.	98.8	100.0	<b>100.0</b>
MultiRC	EM/F1 <sub>a</sub>	65.7 / 88.7	66.3 / 89.0	<b>67.1 / 89.4</b>
ReCoRD	EM/F1	92.7 / 93.4	92.9 / 93.5	<b>93.2 / 93.9</b>
RTE	acc.	92.6	<b>93.5</b>	<b>93.5</b>
WiC	acc.	76.2	76.6	<b>77.4</b>
WSC	acc.	95.8	<b>96.2</b>	<b>96.2</b>
SuperGLUE (dev)		90.5	91.0	<b>91.3</b>

Table 3: Performance of a five-prompt ensemble built from a single frozen T5-XXL model exceeds both the average and the best among the five prompts.

# Conclusion

---

- On SuperGLUE benchmark
  - the task performance rivals that of traditional model tuning, with the gap vanishing as model size increases
- On zero-shot domain transfer
  - prompt tuning leads to improved generalization
- Efficiency
  - efficient storage and serving costs -> multi-task serving, prompt ensembling